

Import Libraries

In [1]:

```
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
import statsmodels.api as sm
from statsmodels.graphics.regressionplots import influence_plot
```

1. Import data set

In [2]:


```
car_data = pd.read_csv('ToyotaCorolla.csv')
car_data
```

Out[2]:

	Id	Model	Price	Age_08_04	Mfg_Month	Mfg_Year	KM	Fuel_Type	HP	Met_...
0	1	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13500	23	10	2002	46986	Diesel	90	
1	2	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13750	23	10	2002	72937	Diesel	90	
2	3	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13950	24	9	2002	41711	Diesel	90	
3	4	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	14950	26	7	2002	48000	Diesel	90	
4	5	TOYOTA Corolla 2.0 D4D HATCHB SOL 2/3- Doors	13750	30	3	2002	38500	Diesel	90	
...
1431	1438	TOYOTA Corolla 1.3 16V HATCHB G6 2/3- Doors	7500	69	12	1998	20544	Petrol	86	
1432	1439	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	10845	72	9	1998	19000	Petrol	86	
1433	1440	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	8500	71	10	1998	17016	Petrol	86	

	Id	Model	Price	Age_08_04	Mfg_Month	Mfg_Year	KM	Fuel_Type	HP	Met (
		TOYOTA								
		Corolla 1.3								
		16V								
1434	1441	HATCHB	7250	70	11	1998	16916	Petrol	86	
		LINEA								
		TERRA								
		2/3-...								
		TOYOTA								
		Corolla 1.6								
1435	1442	LB LINEA	6950	76	5	1998	1	Petrol	110	
		TERRA								
		4/5-Doors								

1436 rows × 38 columns



In [3]:

car_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 38 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1436 non-null   int64
1   Model                 1436 non-null   object
2   Price                 1436 non-null   int64
3   Age_08_04             1436 non-null   int64
4   Mfg_Month             1436 non-null   int64
5   Mfg_Year              1436 non-null   int64
6   KM                    1436 non-null   int64
7   Fuel_Type             1436 non-null   object
8   HP                    1436 non-null   int64
9   Met_Color             1436 non-null   int64
10  Color                 1436 non-null   object
11  Automatic              1436 non-null   int64
12  cc                     1436 non-null   int64
13  Doors                 1436 non-null   int64
14  Cylinders              1436 non-null   int64
15  Gears                  1436 non-null   int64
16  Quarterly_Tax         1436 non-null   int64
17  Weight                 1436 non-null   int64
18  Mfr_Guarantee          1436 non-null   int64
19  BOVAG_Guarantee        1436 non-null   int64
20  Guarantee_Period       1436 non-null   int64
21  ABS                    1436 non-null   int64
22  Airbag_1               1436 non-null   int64
23  Airbag_2               1436 non-null   int64
24  Airco                  1436 non-null   int64
25  Automatic_airco        1436 non-null   int64
26  Boardcomputer          1436 non-null   int64
27  CD_Player              1436 non-null   int64
28  Central_Lock           1436 non-null   int64
29  Powered_Windows        1436 non-null   int64
30  Power_Steering         1436 non-null   int64
31  Radio                  1436 non-null   int64
32  Mistlamps              1436 non-null   int64
33  Sport_Model            1436 non-null   int64
34  Backseat_Divider       1436 non-null   int64
35  Metallic_Rim           1436 non-null   int64
36  Radio_cassette         1436 non-null   int64
37  Tow_Bar                1436 non-null   int64
dtypes: int64(35), object(3)
memory usage: 426.4+ KB
```

In [4]:

```
car_data=pd.DataFrame(data=car_data,columns=["Price","Age_08_04","KM","HP","cc","Doors","Ge
car_data
```

Out[4]:

	Price	Age_08_04	KM	HP	cc	Doors	Gears	Quarterly_Tax	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1431	7500	69	20544	86	1300	3	5	69	1025
1432	10845	72	19000	86	1300	3	5	69	1015
1433	8500	71	17016	86	1300	3	5	69	1015
1434	7250	70	16916	86	1300	3	5	69	1015
1435	6950	76	1	110	1600	5	5	19	1114

1436 rows × 9 columns

In [5]:

```
car_data.isna().sum()
```

Out[5]:

```
Price      0
Age_08_04  0
KM         0
HP         0
cc         0
Doors      0
Gears      0
Quarterly_Tax  0
Weight     0
dtype: int64
```

In [6]:

```
car_data.dtypes
```

Out[6]:

```
Price          int64
Age_08_04      int64
KM             int64
HP             int64
cc             int64
Doors          int64
Gears          int64
Quarterly_Tax  int64
Weight         int64
dtype: object
```

In [7]:

```
car_data = car_data.rename({'Age_08_04':'Age','cc':'CC','Quarterly_Tax':'QT'},axis=1)
car_data
```

Out[7]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1431	7500	69	20544	86	1300	3	5	69	1025
1432	10845	72	19000	86	1300	3	5	69	1015
1433	8500	71	17016	86	1300	3	5	69	1015
1434	7250	70	16916	86	1300	3	5	69	1015
1435	6950	76	1	110	1600	5	5	19	1114

1436 rows × 9 columns

In [8]:

```
car_data.head()
```

Out[8]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170

In [9]:

```
car_data.shape
```

Out[9]:

(1436, 9)

In [10]:

```
car_data.describe(include = 'all')
```

Out[10]:

	Price	Age	KM	HP	CC	Doors	
count	1436.000000	1436.000000	1436.000000	1436.000000	1436.000000	1436.000000	1436.
mean	10730.824513	55.947075	68533.259749	101.502089	1576.85585	4.033426	5.
std	3626.964585	18.599988	37506.448872	14.981080	424.38677	0.952677	0.
min	4350.000000	1.000000	1.000000	69.000000	1300.00000	2.000000	3.
25%	8450.000000	44.000000	43000.000000	90.000000	1400.00000	3.000000	5.
50%	9900.000000	61.000000	63389.500000	110.000000	1600.00000	4.000000	5.
75%	11950.000000	70.000000	87020.750000	110.000000	1600.00000	5.000000	5.
max	32500.000000	80.000000	243000.000000	192.000000	16000.00000	5.000000	6.



In [11]:

```
car_data.corr()
```

Out[11]:

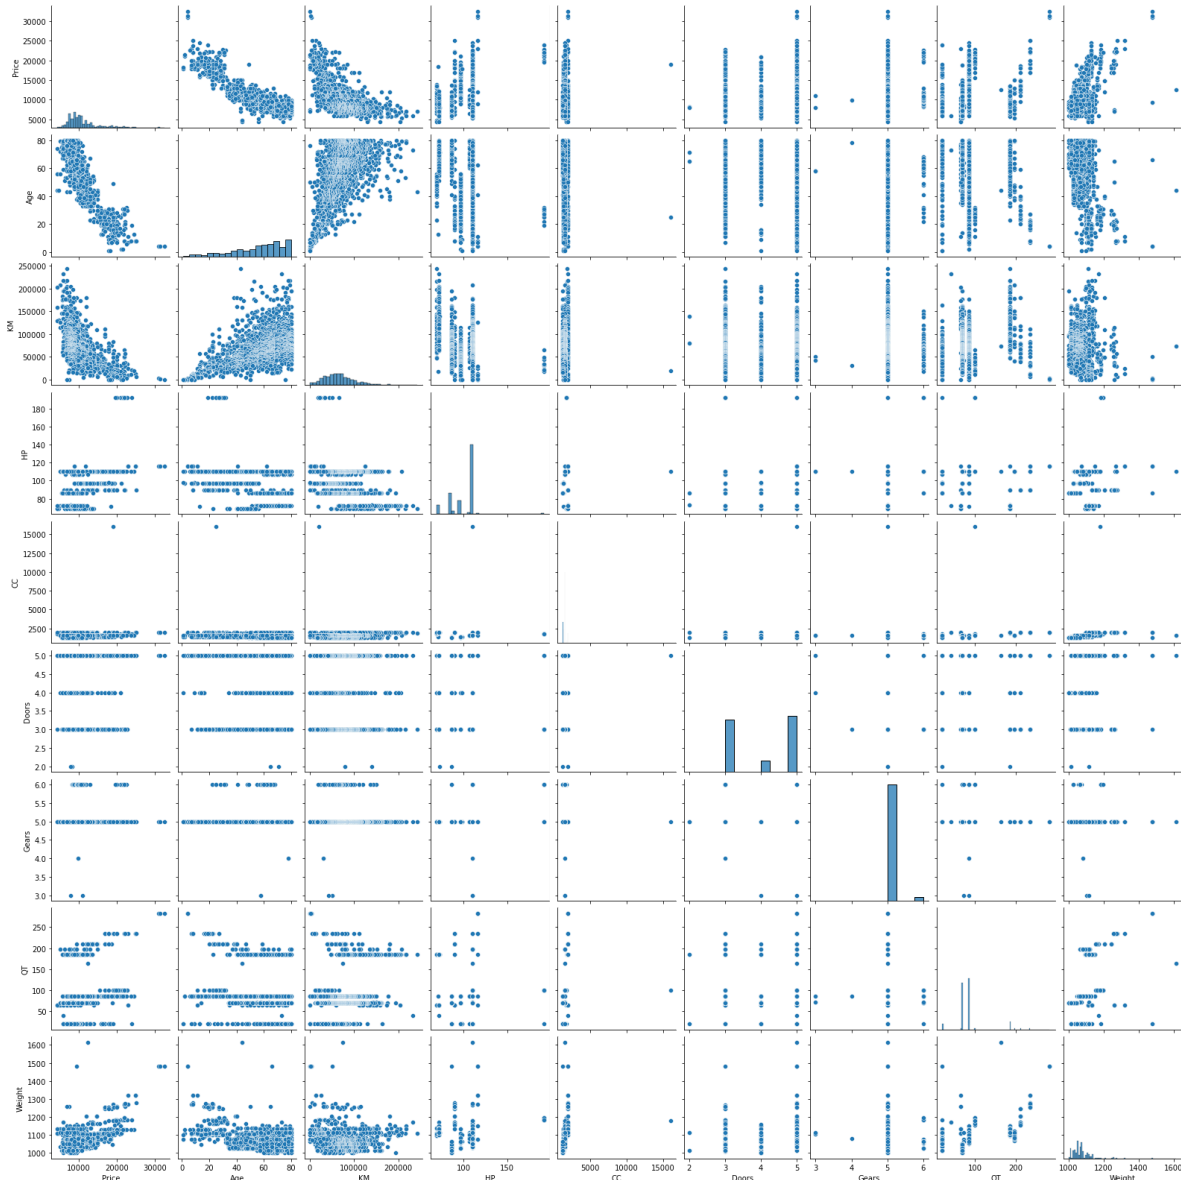
	Price	Age	KM	HP	CC	Doors	Gears	QT
Price	1.000000	-0.876590	-0.569960	0.314990	0.126389	0.185326	0.063104	0.219197
Age	-0.876590	1.000000	0.505672	-0.156622	-0.098084	-0.148359	-0.005364	-0.198431
KM	-0.569960	0.505672	1.000000	-0.333538	0.102683	-0.036197	0.015023	0.278165
HP	0.314990	-0.156622	-0.333538	1.000000	0.035856	0.092424	0.209477	-0.298432
CC	0.126389	-0.098084	0.102683	0.035856	1.000000	0.079903	0.014629	0.306996
Doors	0.185326	-0.148359	-0.036197	0.092424	0.079903	1.000000	-0.160141	0.109363
Gears	0.063104	-0.005364	0.015023	0.209477	0.014629	-0.160141	1.000000	-0.005452
QT	0.219197	-0.198431	0.278165	-0.298432	0.306996	0.109363	-0.005452	1.000000
Weight	0.581198	-0.470253	-0.028598	0.089614	0.335637	0.302618	0.020613	0.626134



2. check for Linearity


In [12]:

```
sns.pairplot(car_data)
plt.show()
```




In [13]:

```
sns.lmplot(x='Age', y='Price', data=car_data)  
plt.show()
```



In [14]:

```
sns.scatterplot( x='Age',y='Price',data=car_data)  
plt.show()
```



3. Model Building

In [15]:

```
model = smf.ols("Price~Age+KM+HP+CC+Doors+Gears+QT+Weight", data=car_data).fit()  
model
```

Out[15]:

```
<statsmodels.regression.linear_model.RegressionResultsWrapper at 0x222aacfba  
30>
```

4. Model Testing

In [16]:

```
# finding p and t values  
np.round(model.pvalues,5),model.tvalues
```

Out[16]:

```
(Intercept      0.00008  
Age             0.00000  
KM              0.00000  
HP              0.00000  
CC              0.17909  
Doors           0.96777  
Gears           0.00261  
QT              0.00262  
Weight          0.00000  
dtype: float64,  
Intercept      -3.948666  
Age            -46.511852  
KM             -16.621622  
HP             11.241018  
CC             -1.344222  
Doors          -0.040410  
Gears           3.016007  
QT              3.014535  
Weight         15.879803  
dtype: float64)
```

In [17]:

```
model.rsquared, model.rsquared_adj
```

Out[17]:

```
(0.8637627463428192, 0.8629989775766963)
```

In [18]:

```
model_2 = smf.ols('Price~CC', data=car_data).fit()  
np.round(model_2.pvalues), model_2.tvalues # CC has Significant pvalue
```

Out[18]:

```
(Intercept    0.0  
CC            0.0  
dtype: float64,  
Intercept    24.694090  
CC           4.824822  
dtype: float64)
```

In [19]:

```
model_3 = smf.ols('Price~Doors', data=car_data).fit()  
model_3
```

Out[19]:

```
<statsmodels.regression.linear_model.RegressionResultsWrapper at 0x222ab6ff490>
```

In [20]:

```
model_3.pvalues,model_3.tvalues # Doors has Significant pvalue
```

Out[20]:

```
(Intercept    1.094732e-73  
Doors         1.461237e-12  
dtype: float64,  
Intercept    19.258097  
Doors         7.141657  
dtype: float64)
```

In [21]:

```
model_4 = smf.ols('Price~CC+Doors', data=car_data).fit()  
model_4.pvalues, model_4.tvalues #CC and Doors have significant pvalues
```

Out[21]:

```
(Intercept    1.056885e-34  
CC            1.521992e-05  
Doors         1.373469e-11  
dtype: float64,  
Intercept    12.620704  
CC            4.340400  
Doors         6.816153  
dtype: float64)
```

Model Validation

In [22]:

```
# Collinearity Check
rsq_age = smf.ols('Age~KM+HP+CC+Doors+Gears+QT+Weight',data=car_data).fit().rsquared
vif_age=1/(1-rsq_age)

rsq_km = smf.ols('KM~Age+HP+CC+Doors+Gears+QT+Weight',data=car_data).fit().rsquared
vif_km=1/(1-rsq_km)

rsq_hp = smf.ols('HP~KM+Age+CC+Doors+Gears+QT+Weight',data=car_data).fit().rsquared
vif_hp=1/(1-rsq_hp)

rsq_cc = smf.ols('CC~HP+KM+Age+Doors+Gears+QT+Weight', data=car_data).fit().rsquared
vif_cc=1/(1-rsq_cc)

rsq_doors = smf.ols('Doors~CC+HP+KM+Age+Gears+QT+Weight',data=car_data).fit().rsquared
vif_doors=1/(1-rsq_doors)

rsq_gears = smf.ols('Gears~Doors+CC+HP+KM+Age+QT+Weight',data=car_data).fit().rsquared
vif_gears=1/(1-rsq_gears)

rsq_qt = smf.ols('QT~Gears+Doors+CC+HP+KM+Age+Weight', data=car_data).fit().rsquared
vif_qt=1/(1-rsq_qt)

rsq_weight = smf.ols('Weight~QT+Gears+Doors+CC+HP+KM+Age',data=car_data).fit().rsquared
vif_weight=1/(1-rsq_weight)


df={'Variables': ['Age', 'KM', 'HP', 'CC', 'Doors', 'Gears', 'QT', 'Weight'],
    'vif': [vif_age, vif_cc, vif_doors, vif_gears, vif_hp, vif_km, vif_qt, vif_weight,]}
vif=pd.DataFrame(df)
vif
```

Out[22]:

	Variables	vif
0	Age	1.884620
1	KM	1.163894
2	HP	1.156575
3	CC	1.098723
4	Doors	1.419422
5	Gears	1.756905
6	QT	2.311431
7	Weight	2.516420

In [23]:

```
# Residual Analysis
sm.qqplot(model.resid, line='q')
plt.title('Noarmal Q-Q plots of residuals')
plt.show()
```



In [24]:

```
list(np.where(model.resid>6000))
```

Out[24]:

```
[array([147, 523], dtype=int64)]
```

In [25]:


```
list(np.where(model.resid<-6000))
```

Out[25]:

```
[array([221, 601, 960], dtype=int64)]
```

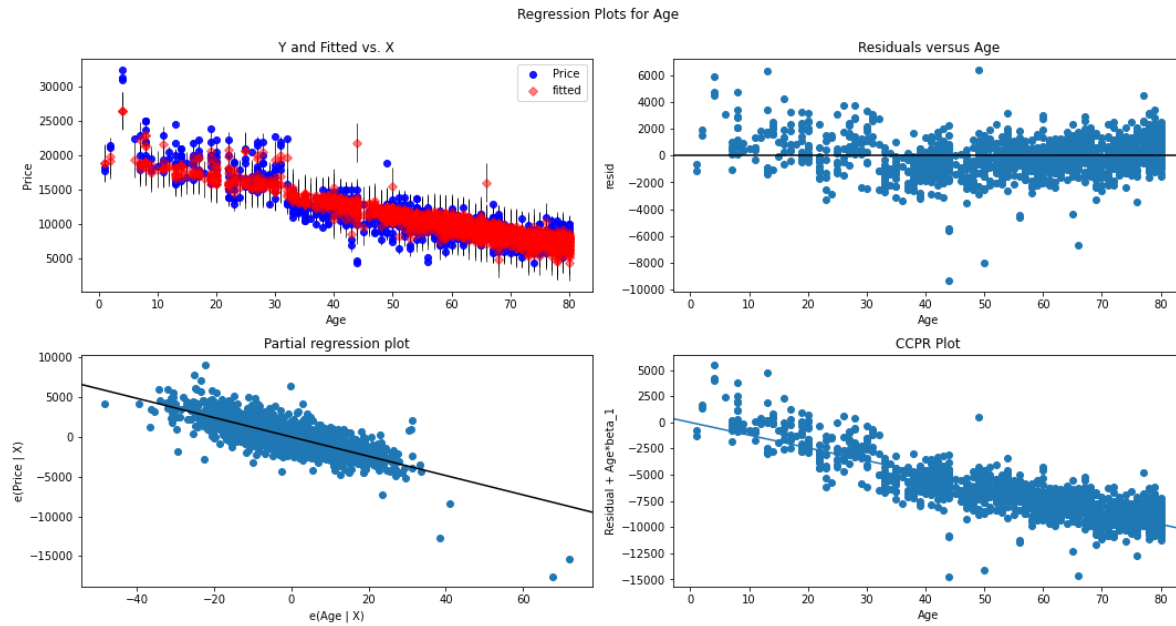
In [26]:

```
def standard_values(vals) : return (vals-vals.mean())/vals.std()  
plt.scatter(standard_values(model.fittedvalues),standard_values(model.resid))  
plt.title('Residual plot')  
plt.xlabel('standardized fitted values')  
plt.ylabel('standardized residual values')  
plt.show()
```



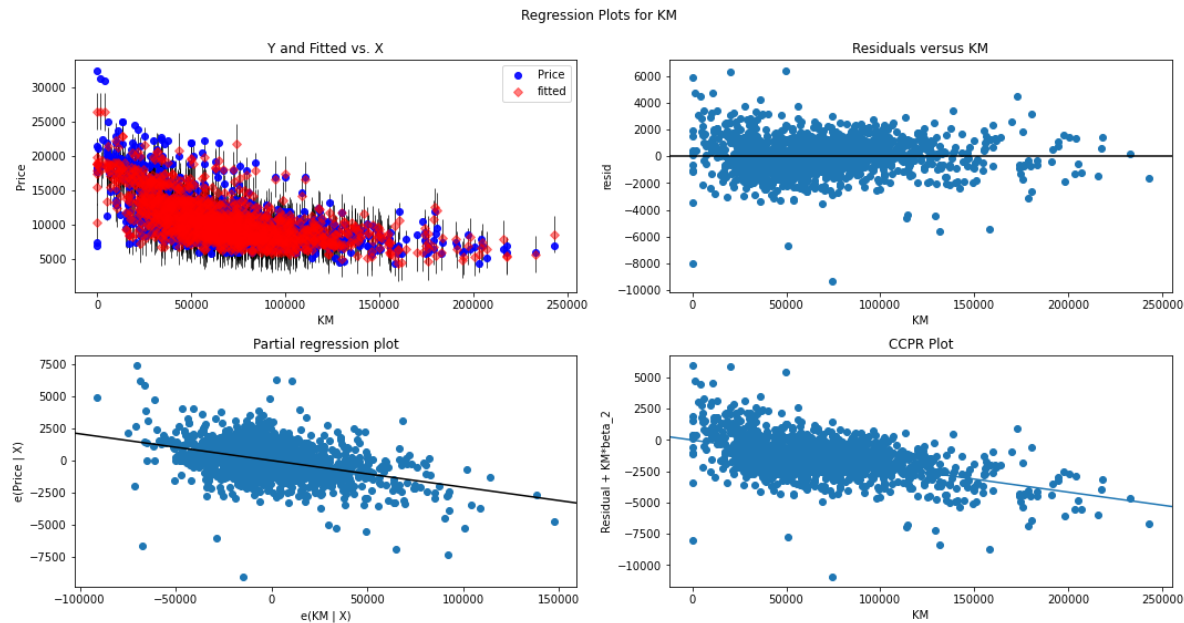
In [27]:

```
#residual plots
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Age',fig=fig)
plt.show()
```



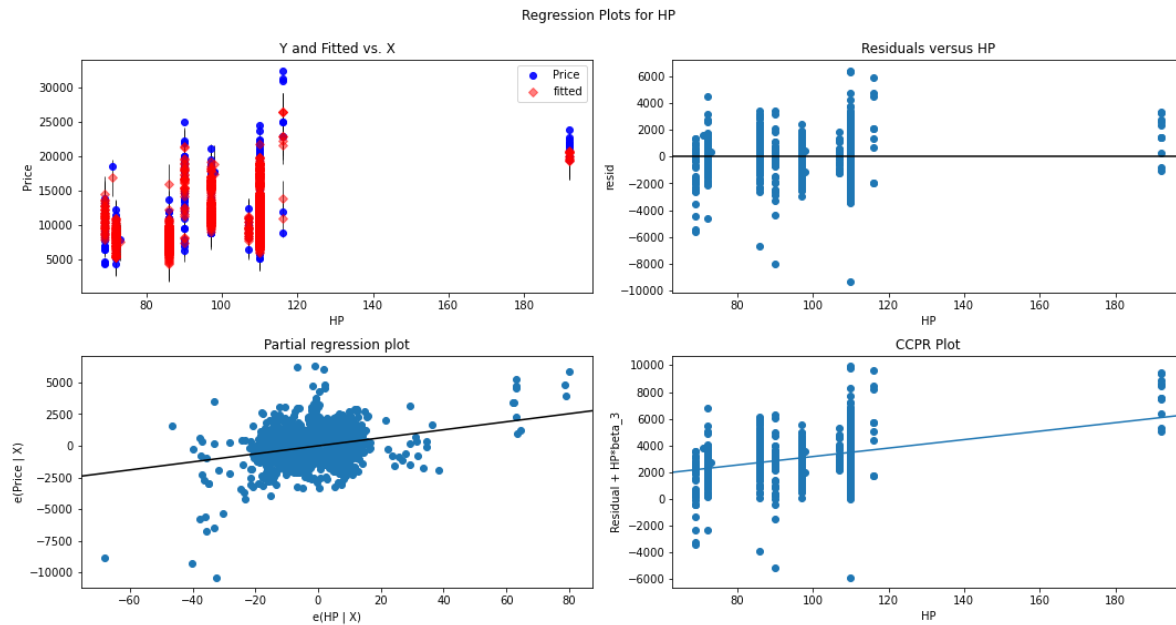
In [28]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'KM',fig=fig)
plt.show()
```



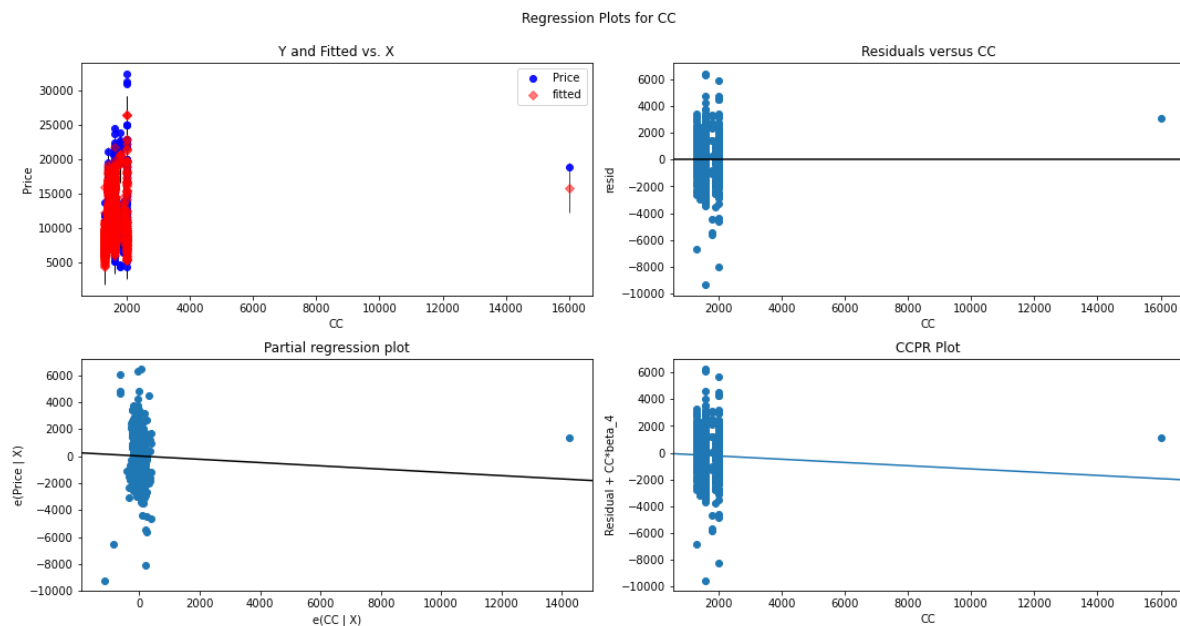
In [29]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'HP',fig=fig)
plt.show()
```



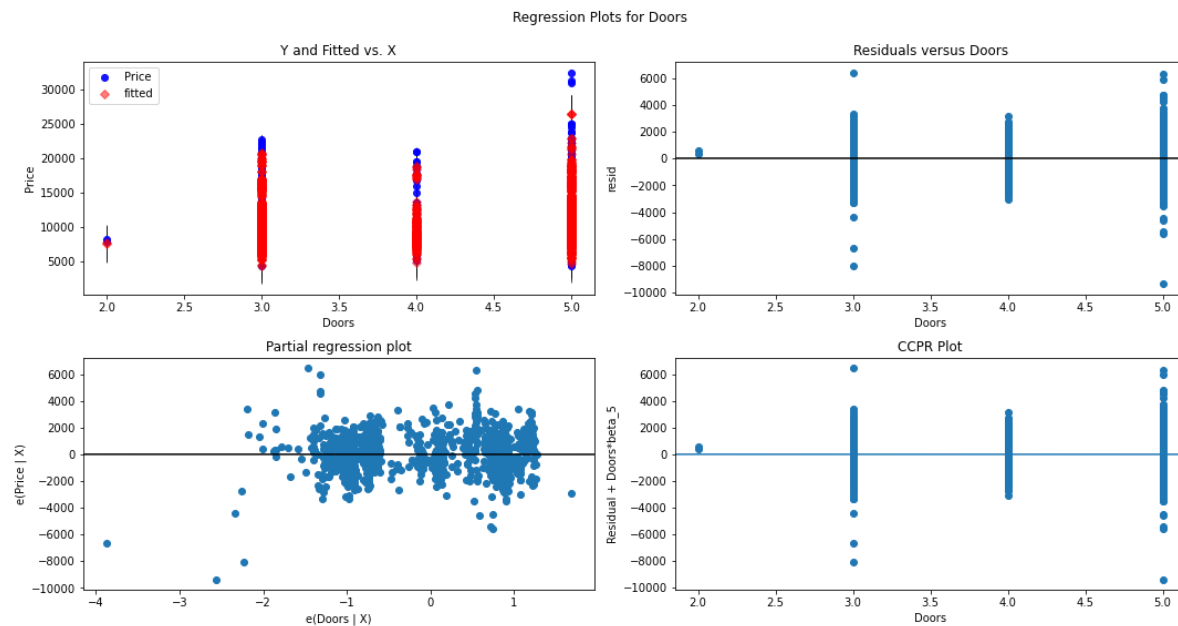
In [30]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'CC',fig=fig)
plt.show()
```




In [31]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Doors',fig=fig)
plt.show()
```



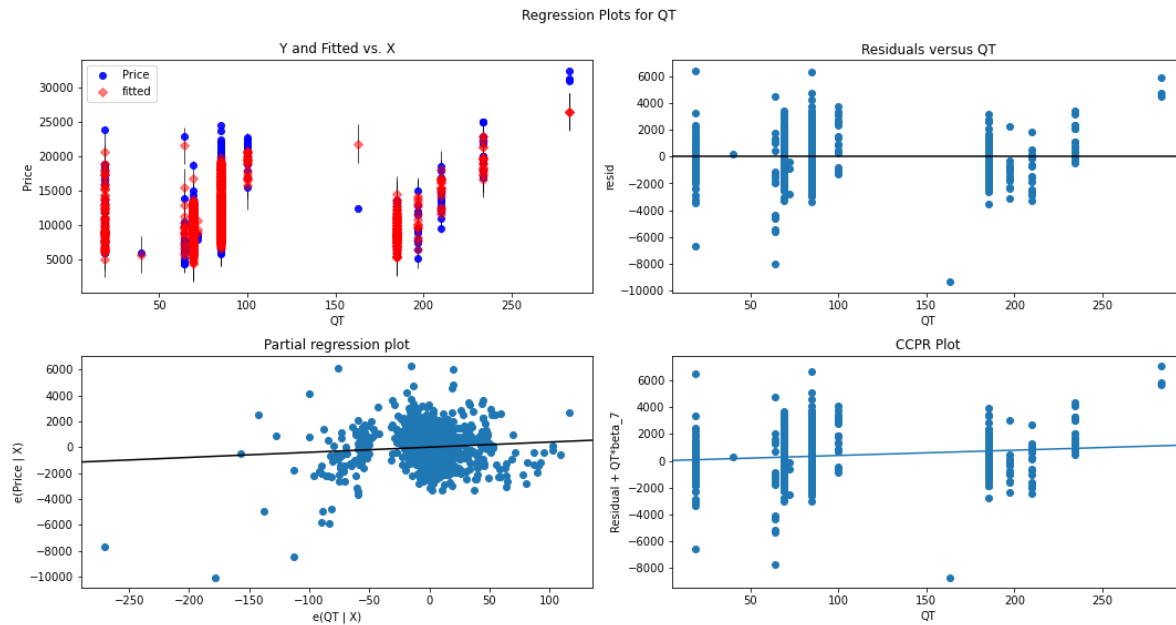
In [32]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Gears',fig=fig)
plt.show()
```



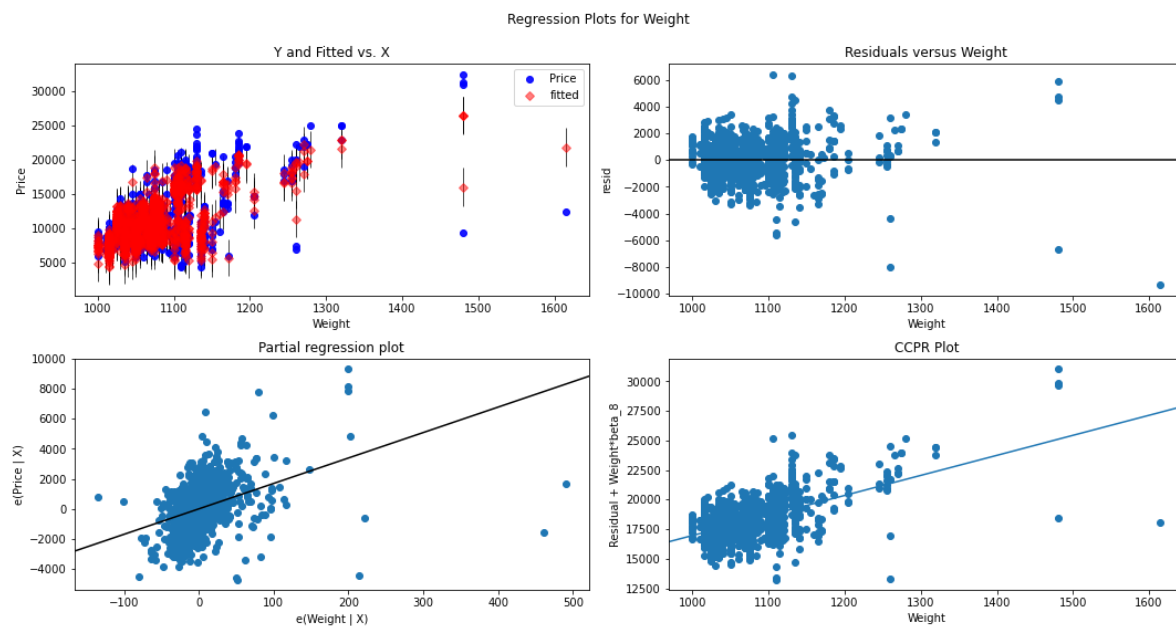
In [33]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'QT',fig=fig)
plt.show()
```



In [34]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'Weight',fig=fig)
plt.show()
```



In [35]:


```
# Model delegation diagogics
#1. cooks distence
(C,_)=model.get_influence().cooks_distance
C
```

Out[35]:

```
array([7.23682667e-03, 3.96793393e-03, 5.46476784e-03, ...,
       8.44762355e-07, 6.97878368e-04, 1.08627724e-02])
```

In [36]:

```
#plot the influencers using stem plot
plt.figure(figsize=(25,8))
plt.stem(np.arange(len(car_data)),np.round(C,3))
plt.xlabel('Row index')
plt.ylabel('Cooks distence')
plt.show()
```




In [37]:

```
np.argmax(C),np.max(C)
```

Out[37]:

```
(80, 79.52010624138181)
```

```
#2.Leverage value using high influence points
fig,ax=plt.subplots(figsize=(20,20))
fig=influence_plot(model, ax = ax)
```



In [39]:

```
# Leverage Cutoff values
k=car_data.shape[1]
n=car_data.shape[0]
levarage_cutoff=(3*(k+1))/n
levarage_cutoff
```

Out[39]:

0.020891364902506964

In [40]:

```
car_data[car_data.index.isin([80])]
```

Out[40]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
80	18950	25	20019	110	16000	5	5	100	1180

In [41]:

```
# improving model
car_new = car_data.copy()
car_new
```

Out[41]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1431	7500	69	20544	86	1300	3	5	69	1025
1432	10845	72	19000	86	1300	3	5	69	1015
1433	8500	71	17016	86	1300	3	5	69	1015
1434	7250	70	16916	86	1300	3	5	69	1015
1435	6950	76	1	110	1600	5	5	19	1114

1436 rows × 9 columns

In [42]:

```
car_data1 = car_new.drop(car_new.index[[80]], axis=0).reset_index()
car_data1
```

Out[42]:

	index	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	0	13500	23	46986	90	2000	3	5	210	1165
1	1	13750	23	72937	90	2000	3	5	210	1165
2	2	13950	24	41711	90	2000	3	5	210	1165
3	3	14950	26	48000	90	2000	3	5	210	1165
4	4	13750	30	38500	90	2000	3	5	210	1170
...
1430	1431	7500	69	20544	86	1300	3	5	69	1025
1431	1432	10845	72	19000	86	1300	3	5	69	1015
1432	1433	8500	71	17016	86	1300	3	5	69	1015
1433	1434	7250	70	16916	86	1300	3	5	69	1015
1434	1435	6950	76	1	110	1600	5	5	19	1114

1435 rows × 10 columns

In [43]:

```
car_data1 = car_data1.drop(['index'], axis=1)
car_data1
```

Out[43]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...
1430	7500	69	20544	86	1300	3	5	69	1025
1431	10845	72	19000	86	1300	3	5	69	1015
1432	8500	71	17016	86	1300	3	5	69	1015
1433	7250	70	16916	86	1300	3	5	69	1015
1434	6950	76	1	110	1600	5	5	19	1114

1435 rows × 9 columns

In [55]:

```
while np.max(C) > 0.5:
    model = smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=car_data1).fit()
    (C,_)=model.get_influence().cooks_distance
    C
    np.argmax(C),np.max(C)
    car_data1 = car_data1.drop(car_data1.index[[np.argmax(C)]], axis=1).reset_index()
    car_data1
else:
    final_model = smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=car_data1).fit()
    final_model.rsquared,final_model.aic
    print('Thus model accuracy is improved to', final_model.rsquared)
```

Thus model accuracy is improved to 0.8851845904421739

In [57]:

```
if np.max(C)>0.5 :
    model = smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=car_data1).fit()
    (C,_)=model.get_influence().cooks_distance
    C
    np.argmax(C),np.max(C)
    car_data1=car_data1.drop(car_data1.index[[np.argmax(C)]],axis=1).reset_index()
    car_data1
elif np.max(C)<0.5:
    final_model=smf.ols('Price~Age+KM+HP+CC+Doors+Gears+QT+Weight',data=car_data1).fit()
    final_model.rsquared,final_model.aic
    print('Thus model accuracy is improved to',final_model.rsquared)
```

Thus model accuracy is improved to 0.8851845904421739

In [58]:

```
final_model.rsquared
```

Out[58]:

0.8851845904421739

In [59]:

```
car_data1
```

Out[59]:

	level_0	index	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	0	0	13500	23	46986	90	2000	3	5	210	1165
1	1	1	13750	23	72937	90	2000	3	5	210	1165
2	2	2	13950	24	41711	90	2000	3	5	210	1165
3	3	3	14950	26	48000	90	2000	3	5	210	1165
4	4	4	13750	30	38500	90	2000	3	5	210	1170
...
1428	1429	1430	7500	69	20544	86	1300	3	5	69	1025
1429	1430	1431	10845	72	19000	86	1300	3	5	69	1015
1430	1431	1432	8500	71	17016	86	1300	3	5	69	1015
1431	1432	1433	7250	70	16916	86	1300	3	5	69	1015
1432	1433	1434	6950	76	1	110	1600	5	5	19	1114

1433 rows × 11 columns

In [70]:

```
# Model prediction for new data
new_data = pd.DataFrame({'Age':15, 'KM':50000, 'HP':90, 'CC':1400, 'Doors':4, 'Gears':5, 'QT':210})
new_data
```

Out[70]:

	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	15	50000	90	1400	4	5	210	1165

In [72]:

```
final_model.predict(new_data)
```

Out[72]:

```
0    19332.917337
dtype: float64
```

In [74]:

```
pred_y = final_model.predict(car_data1)
pred_y
```

Out[74]:

```
0      16333.273814
1      15892.326850
2      16310.886081
3      15979.990390
4      15846.536733
...
1428    9115.435074
1429    8499.218117
1430    8644.947302
1431    8758.664462
1432   10641.521002
Length: 1433, dtype: float64
```

In []: